

Drift

an imperative programming
environment for the cloud

#2 Implementation

Overview


- Drift Language

- Concepts
- Examples

- Drift Execution

- Drift FS
- Architecture
- Error Model

Recap

- Want: ‘language of the system’
- Workflow languages one possible domain
 - black box tasks, ...
- Bash: system-view coordination
 - black box tasks, immediate feedback, ...
 - Problem: FS → shared mutable state
- Functional  Distributed
 - Cuneiform: functional, distributed, ...

Recap

Bash	Functional
&	- (lazy)
	function composition
>, >>	name binding
<	-
\$	eval

Drift Language

Can we build an *imperative, stateful, interpreted* language for distributed (micro) service coordination ?

- Need ‘state’ to be stateful on
- Essence: data + services

“Do this on that thing over here.”

“Now do this on that and put it over there”

Drift Language

- How do we (humans) ‘interact’ with data?
 - need ‘names’ to identify and retrieve our data
 - need ‘names’ to give data *meaning*
- *Names* are at the center of programming!
 - only names and services
- Names best be hierarchical → Namespaces
 - Names, Namespaces and Services

Drift Language

```
.> Cat foo.txt
  Lorem ipsum dolor sit amet, consetetur
  sadipscing elitr, sed diam nonumy eirmod
.>

.> foo = Cat foo.txt
.>

.> ls
  foo

.> $foo
  Lorem ipsum dolor sit amet, consetetur
  sadipscing elitr, sed diam nonumy eirmod
  ...
```


Drift Language

```
.> wordcount_h = Wc hamlet.txt  
.> wordcount_mb = Wc macbeth.txt  
.> Max wordcount_h wordcount_mb  
wordcount_h
```

```
.> a = A data.csv | B | C
```

```
.> a = A in1.data in2.data  
.> b = B a | C  
.> a = G homework.txt
```

Pipe cut into
single commands

Drift Language

```
.> import my.tar  
.> res/ = Untar* my.tar | *FormatCheck txt  
  
.> ls  
  my.tar  
  res/  
  
.> $res/  
  c1.txt  
  c2.txt  
  
.> wordcount = Wc res/c1.txt
```

ls, cd, rm are
language keywords

Drift Language

So far:

- no arithmetic
- no conditionals
- not turing-complete! (hopefully)
- tiny
- very abstract
- very few constructs

Drift Execution

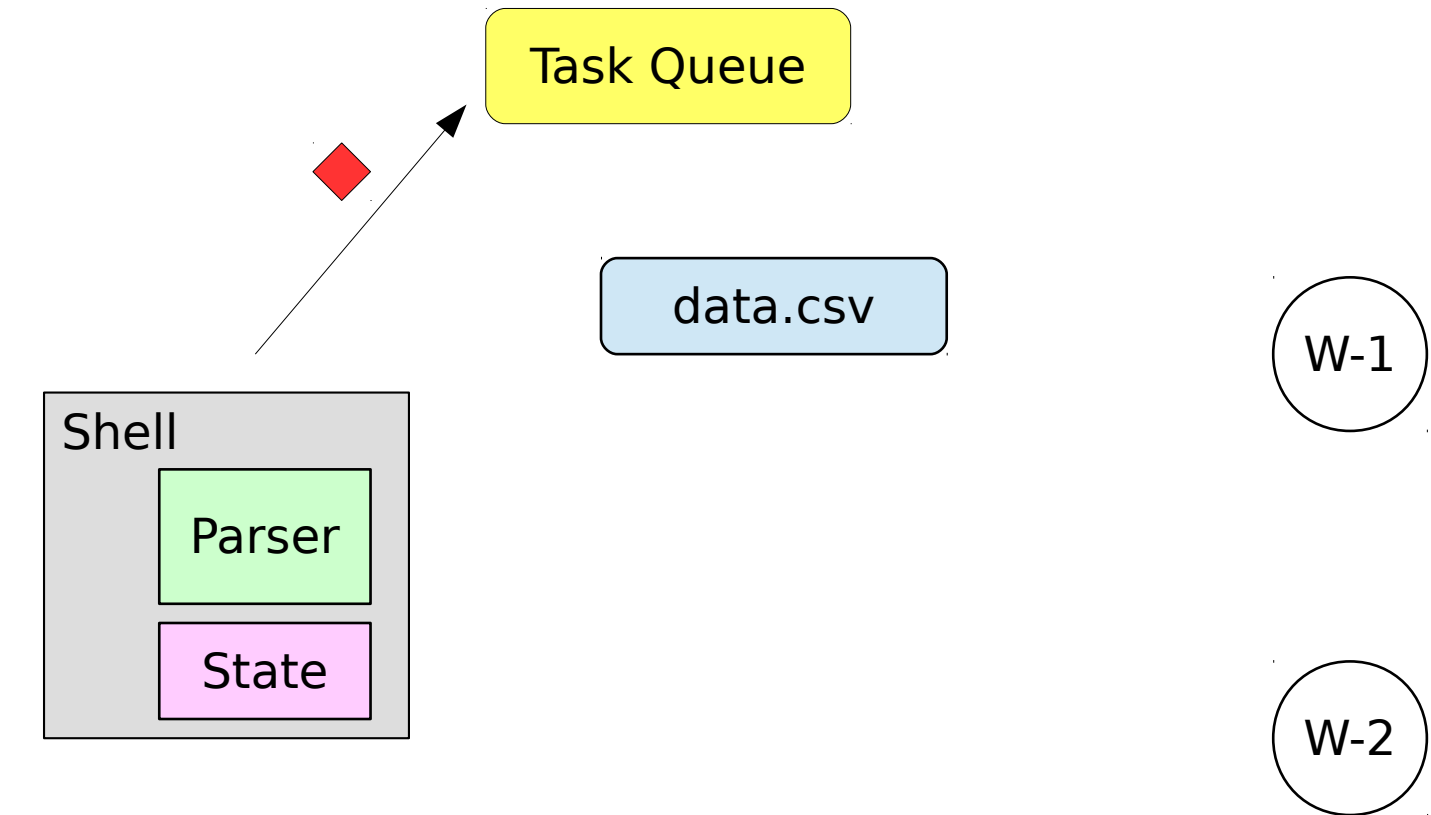
Front End:

- own shell implementation
- ANTLR 4 for language parsing



Back End:

- own worker implementation
- RabbitMQ and Kafka as data and signal queues
- Mesos + Marathon for fault-tolerance (Workers)

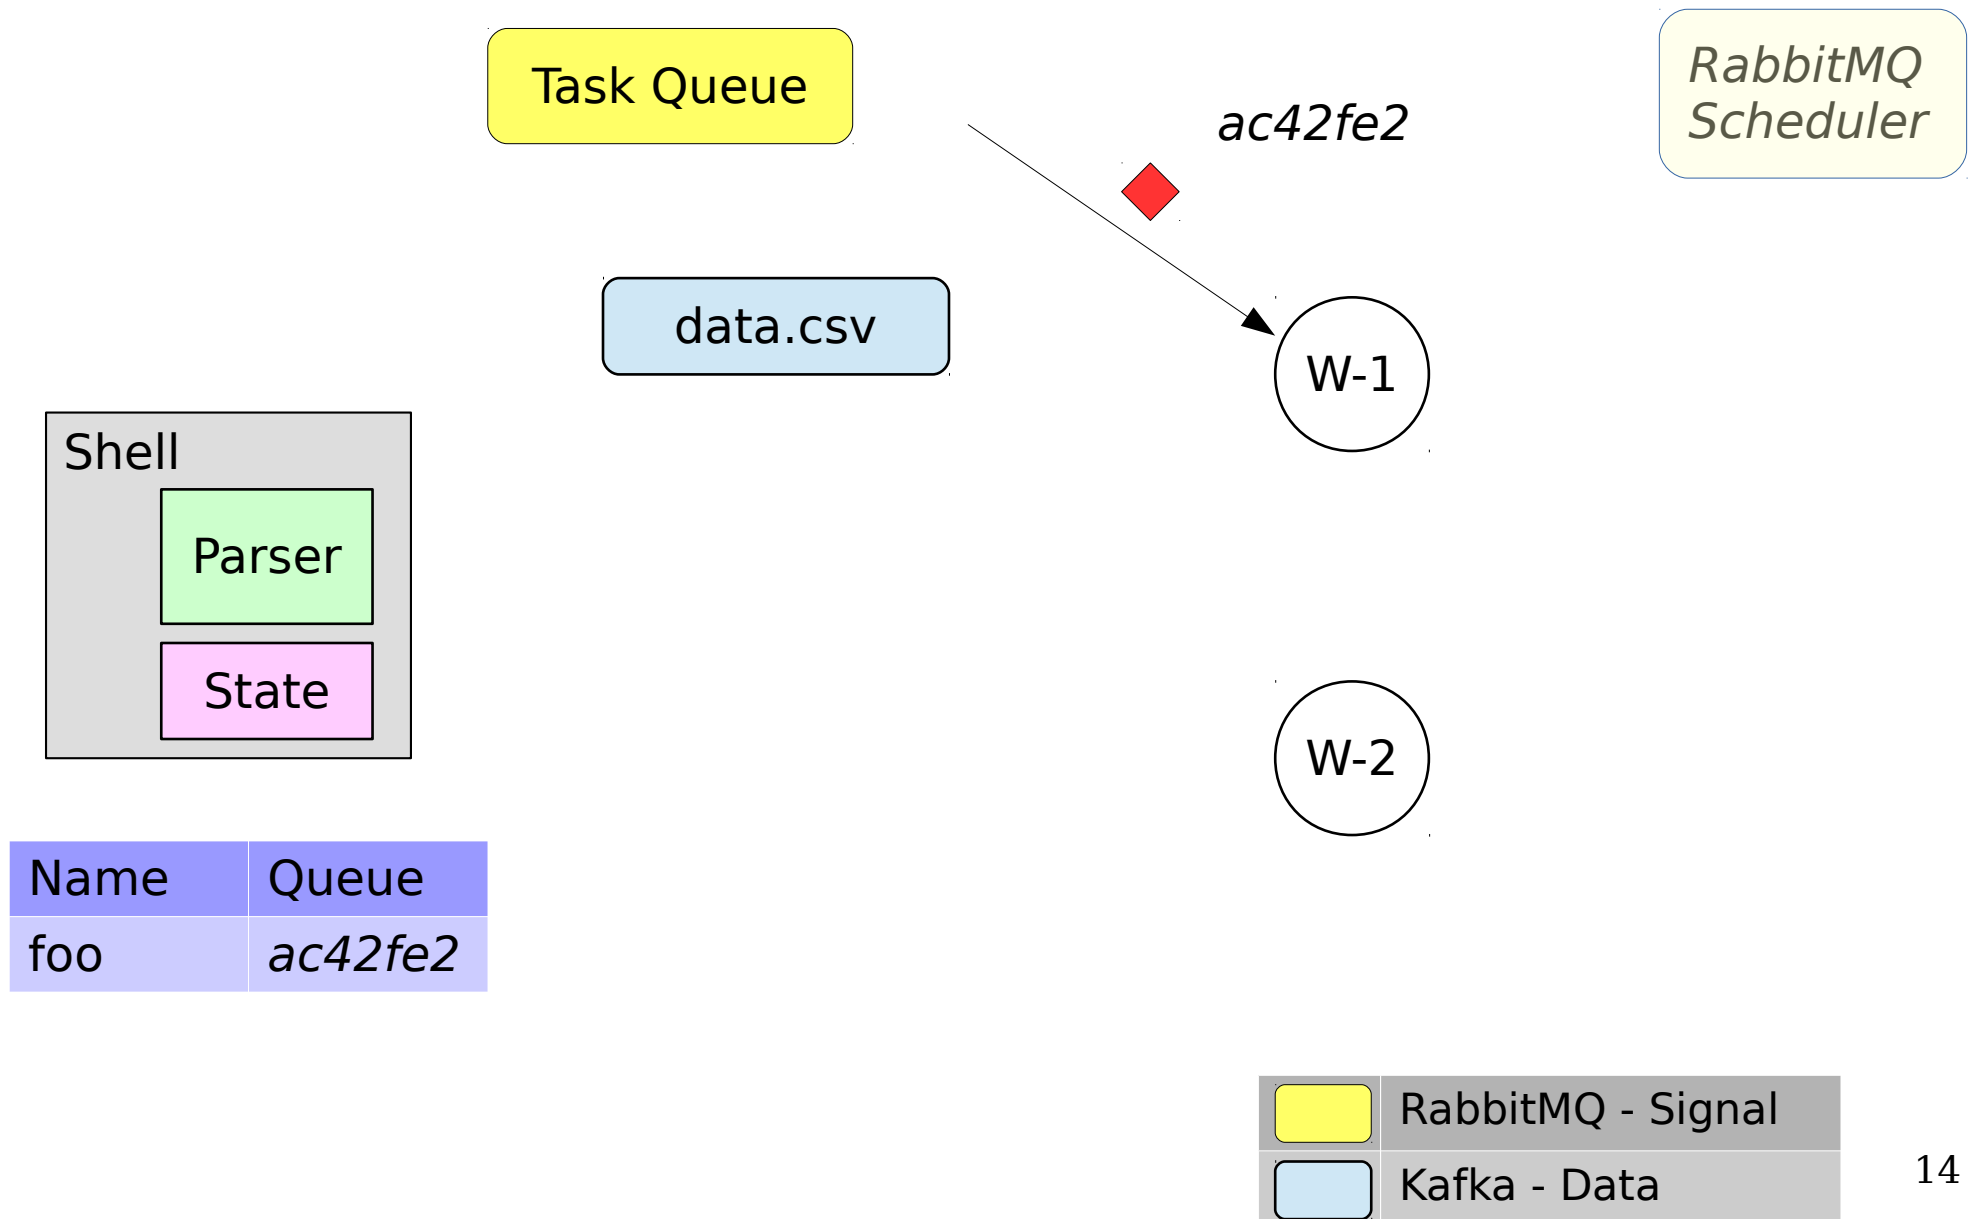
Drift Execution



Name	Queue
foo	<i>ac42fe2</i>

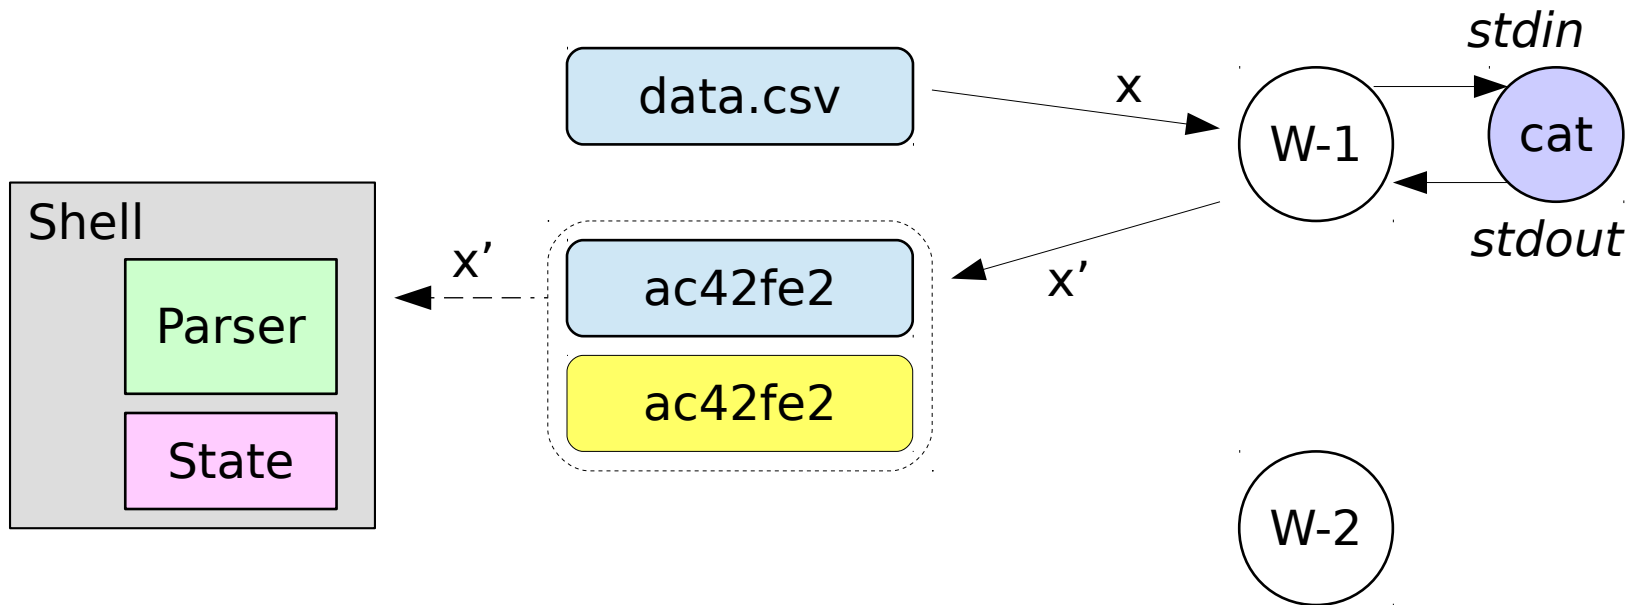
	RabbitMQ - Signal
	Kafka - Data

Drift Execution





Drift Execution

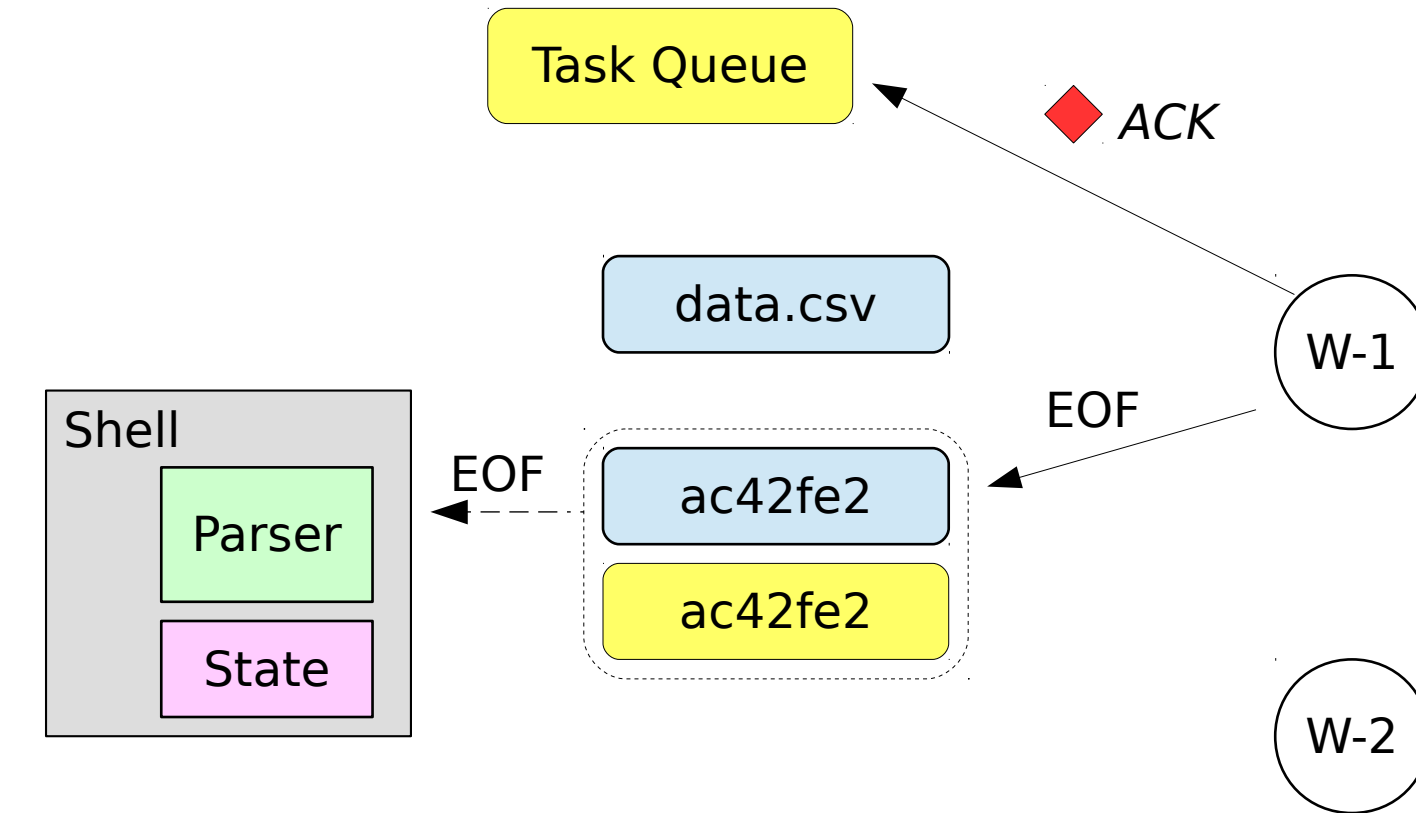
Task Queue





Name	Queue
foo	ac42fe2

	RabbitMQ - Signal
	Kafka - Data

Drift Execution

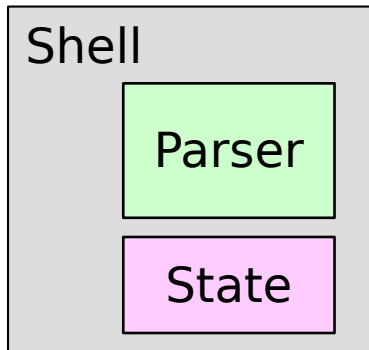


Name	Queue
foo	ac42fe2

	RabbitMQ - Signal
	Kafka - Data

Drift Execution

Task Queue



data.csv



ac42fe2 ✓

ac42fe2

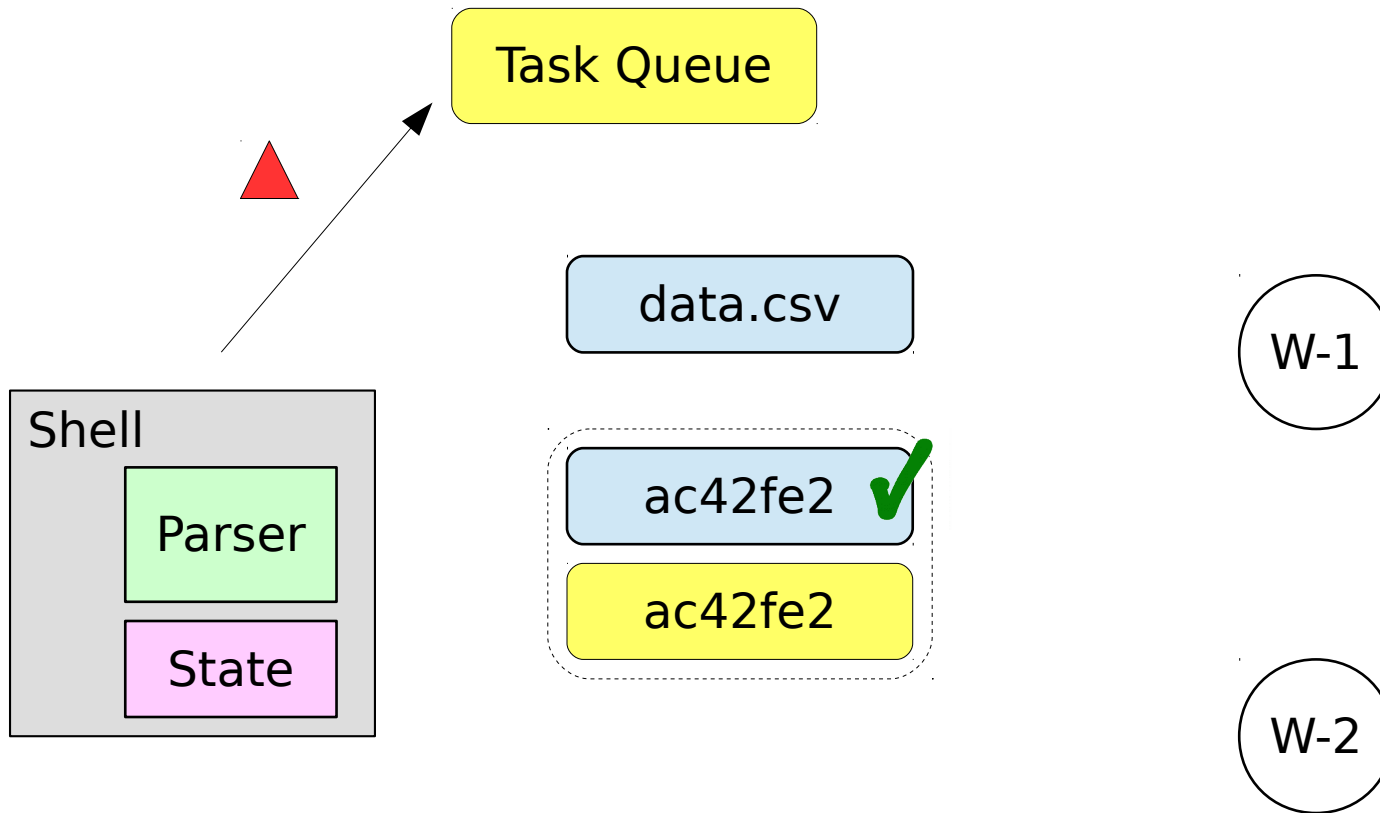
W-1

W-2



Name	Queue
foo	ac42fe2

	RabbitMQ - Signal
	Kafka - Data

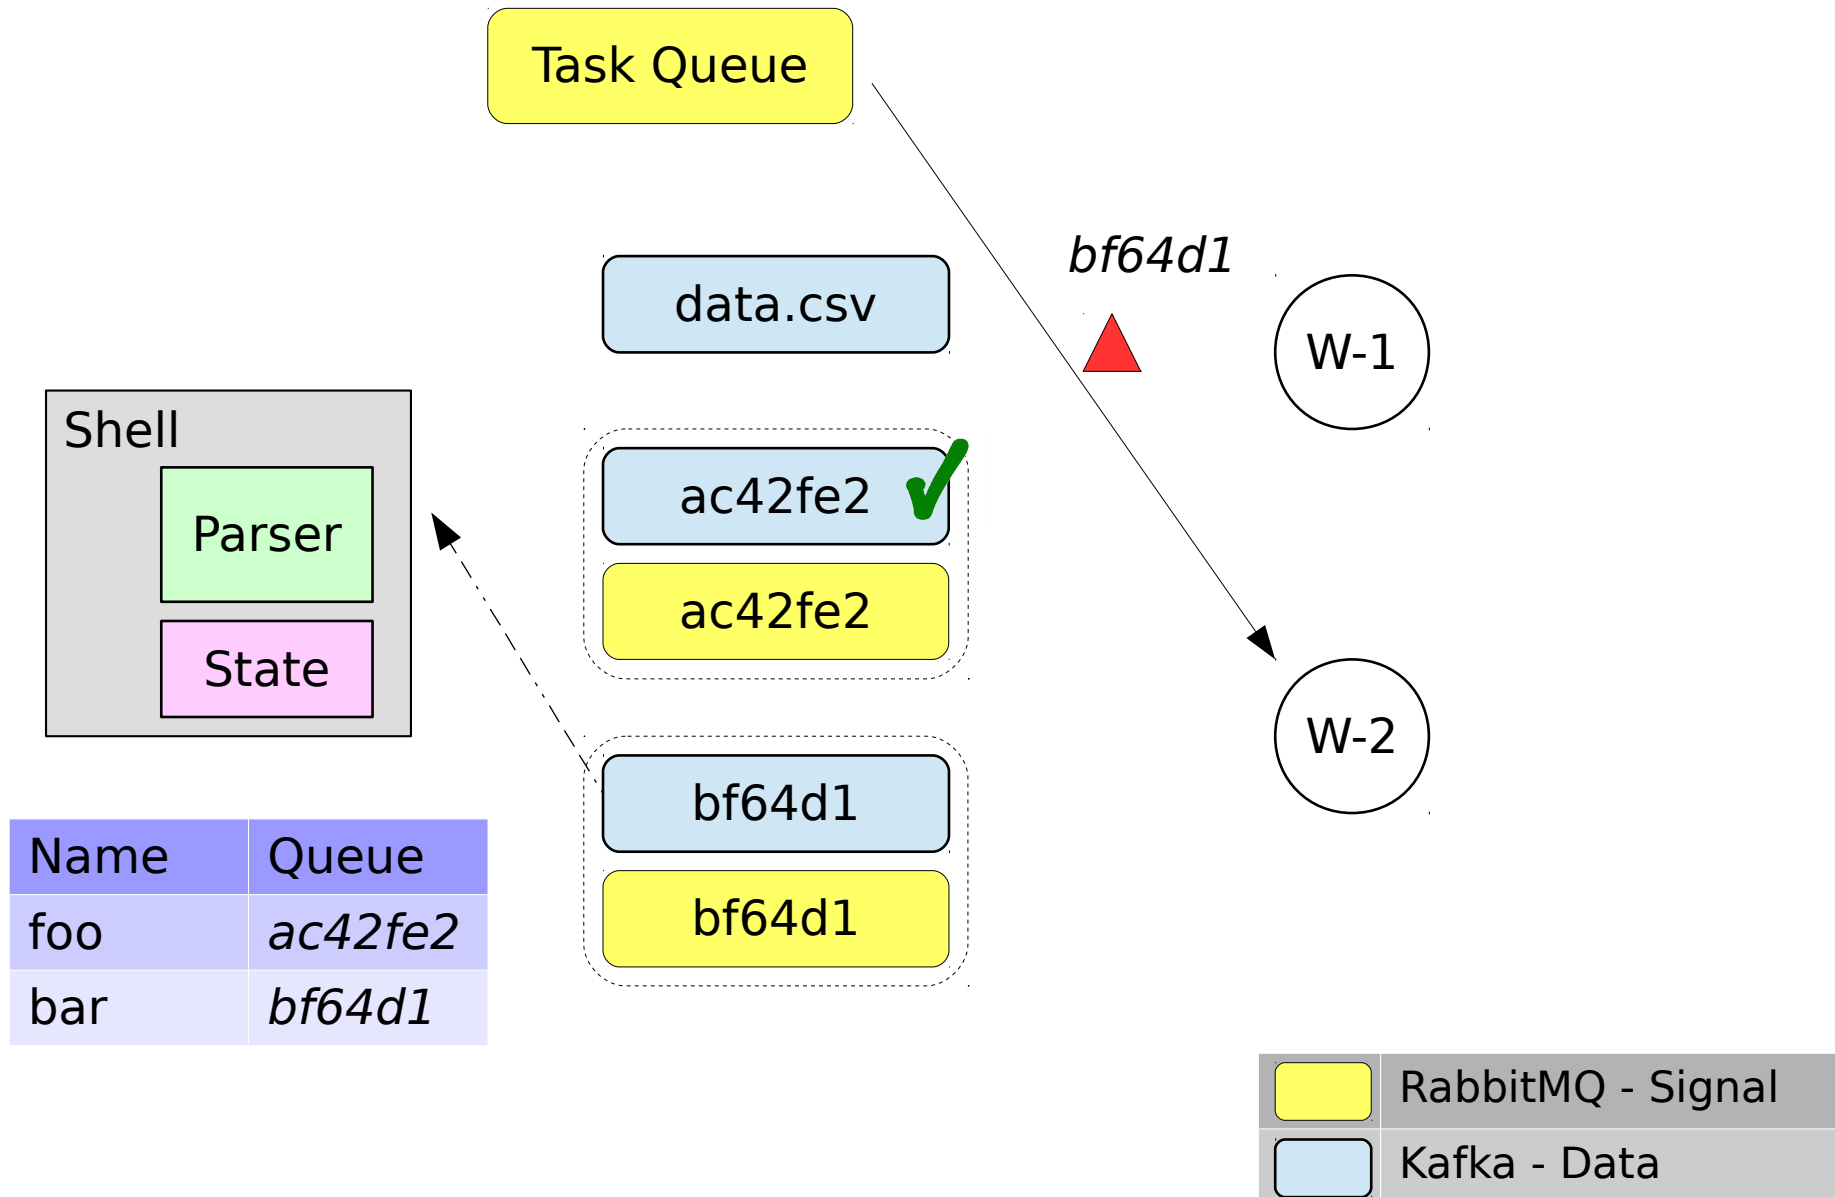
Drift Execution



Name	Queue
foo	<i>ac42fe2</i>
bar	<i>bf64d1</i>

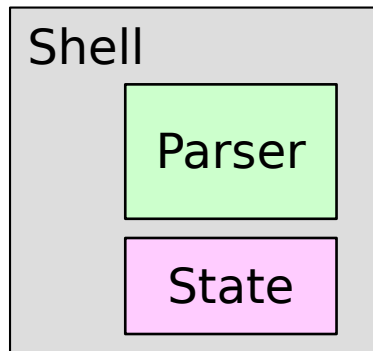
	RabbitMQ - Signal
	Kafka - Data

Drift Execution



Drift Execution

Task Queue



Name	Queue
foo	ac42fe2
bar	bf64d1

data.csv

ac42fe2 ✓

ac42fe2

bf64d1

bf64d1

W-1

y

W-2

y'

stdin

stdout

WC



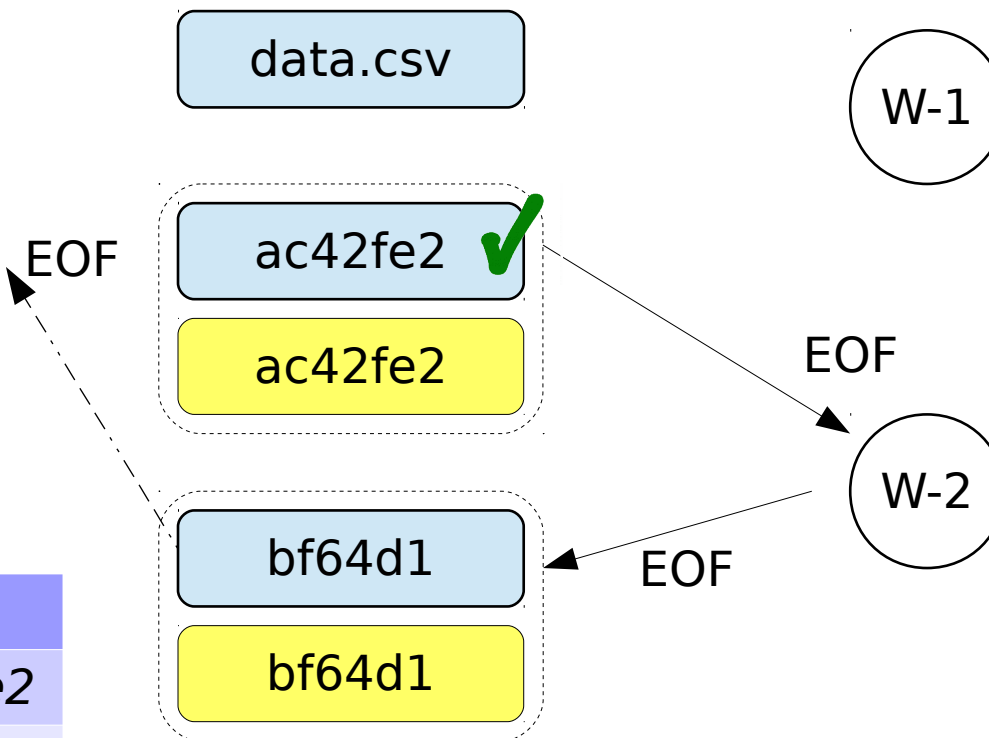
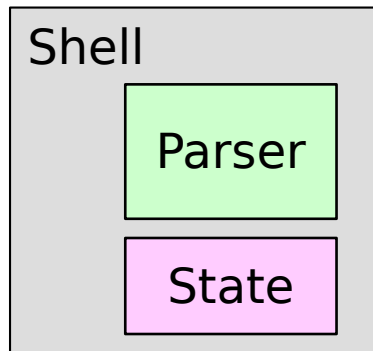
RabbitMQ - Signal





Kafka - Data

Drift Execution

Task Queue

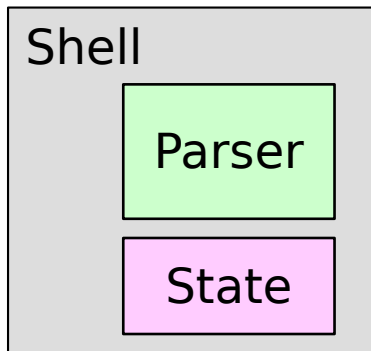


Name	Queue
foo	<i>ac42fe2</i>
bar	<i>bf64d1</i>

	RabbitMQ - Signal
	Kafka - Data

Drift Execution

Task Queue



Name	Queue
foo	<i>ac42fe2</i>
bar	<i>bf64d1</i>

data.csv

ac42fe2 ✓

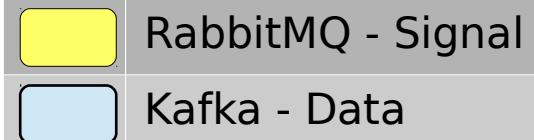
ac42fe2

bf64d1 ✓

bf64d1

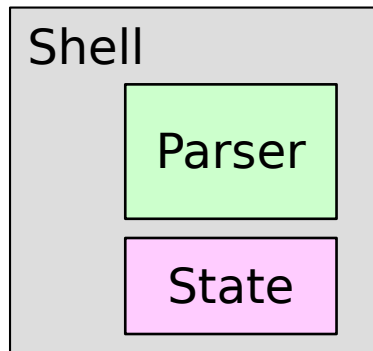
W-1

W-2

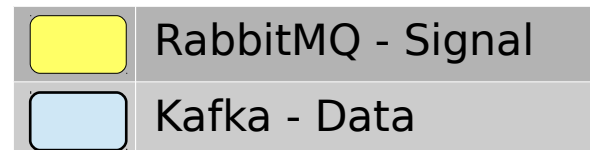
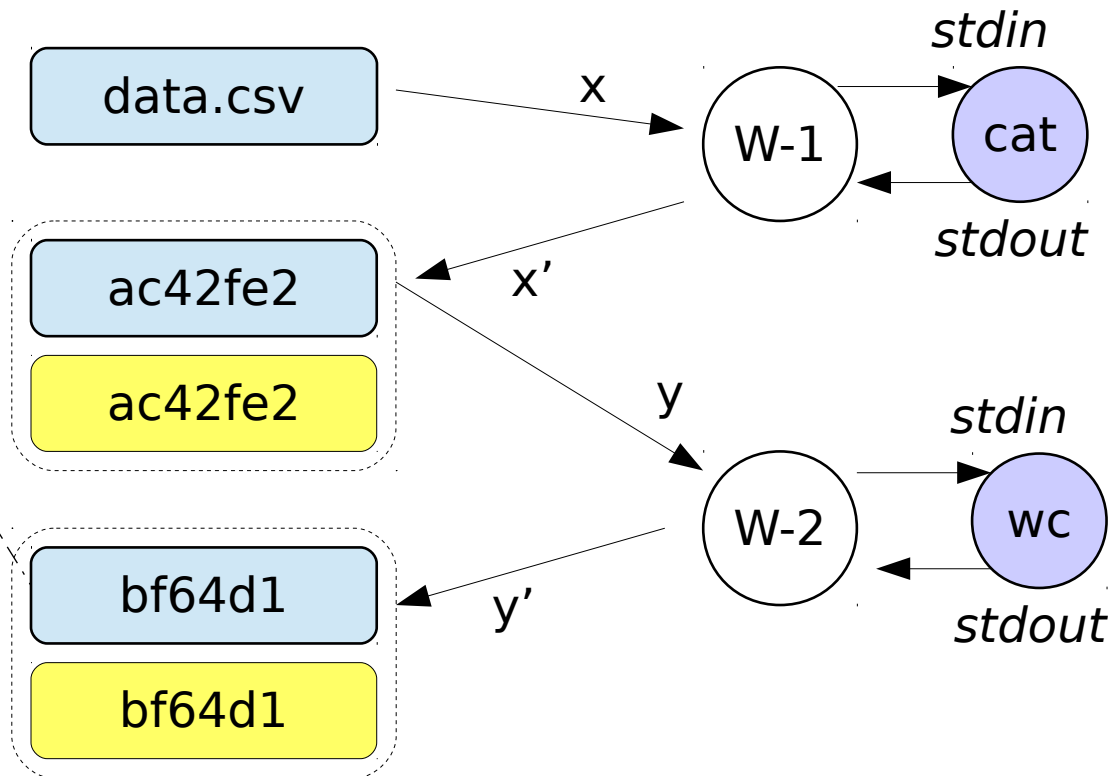


Drift Execution

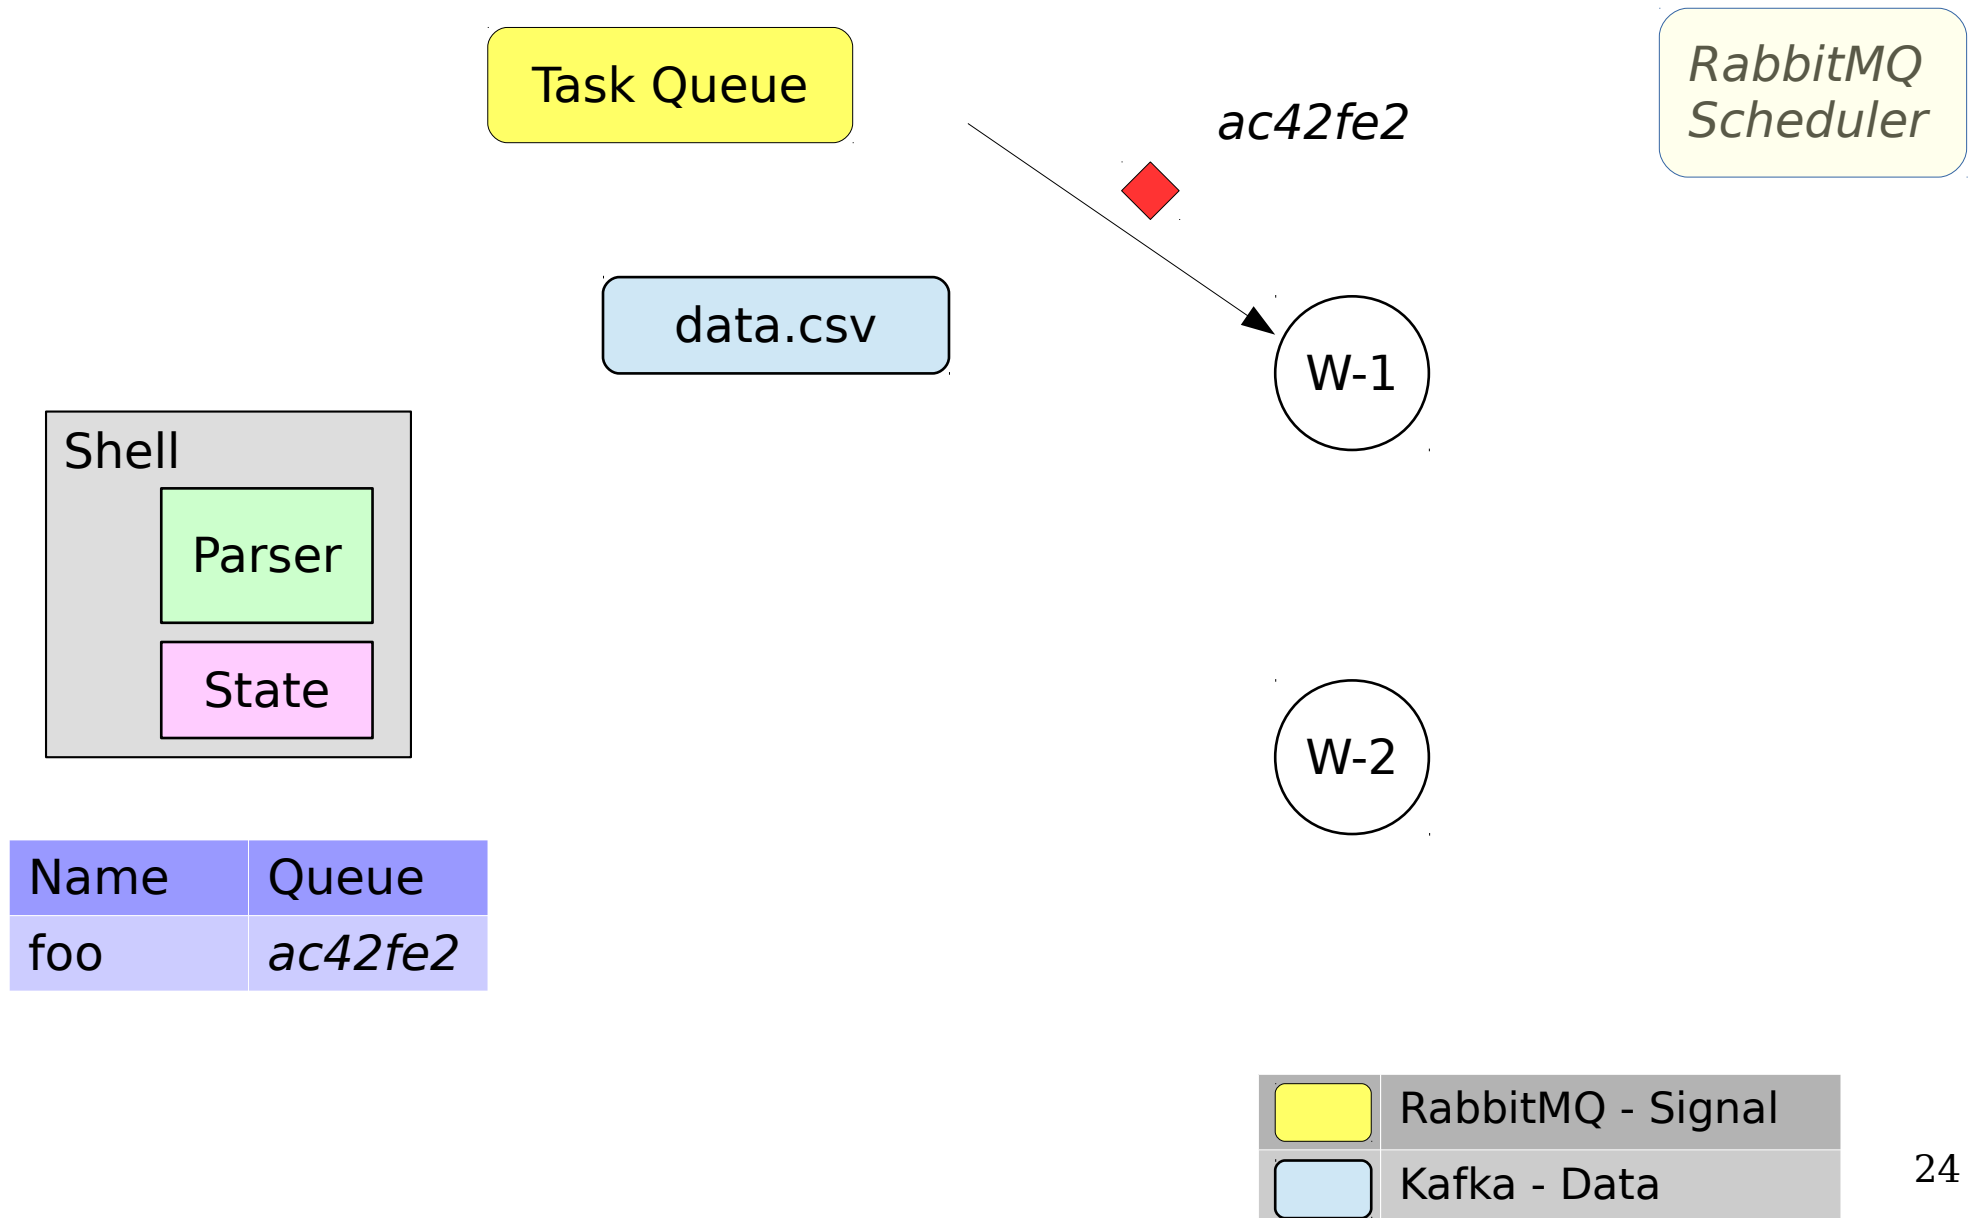
Task Queue



Name	Queue
foo	<i>ac42fe2</i>
bar	<i>bf64d1</i>



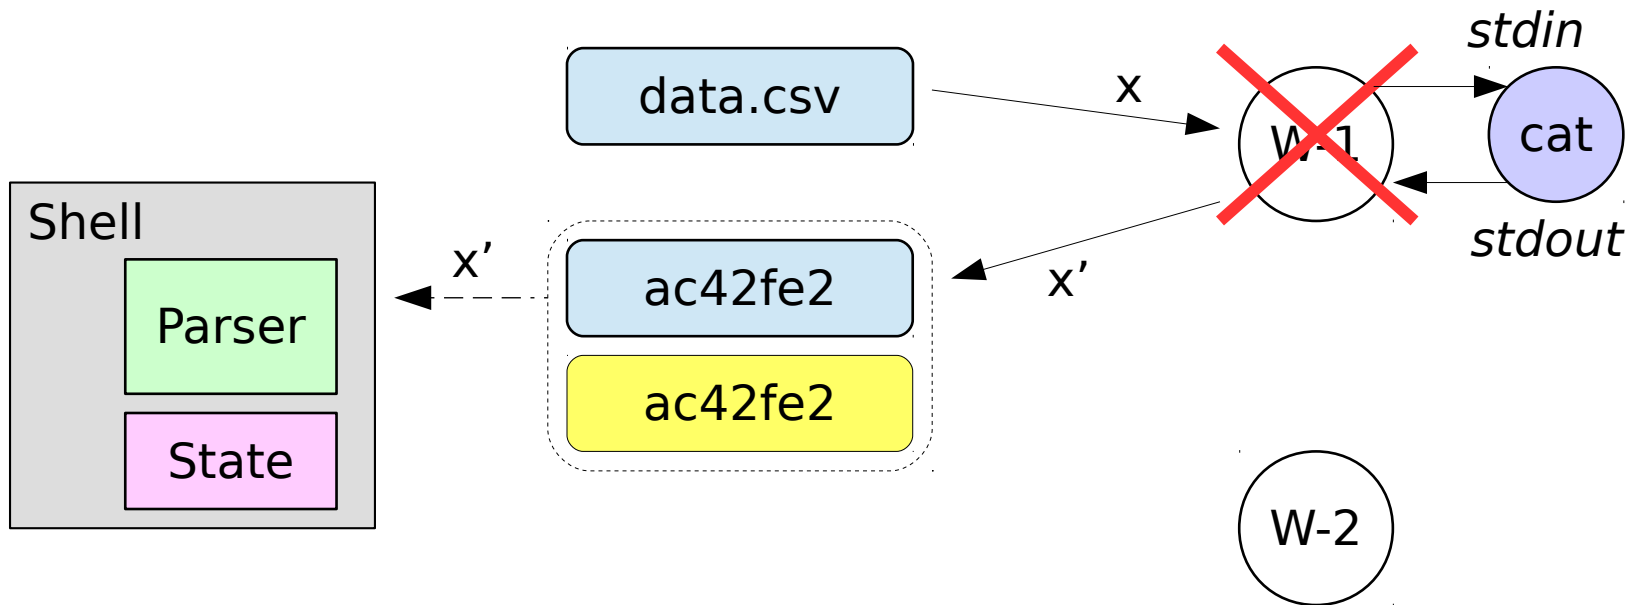
Drift Execution





Drift Execution

Task Queue

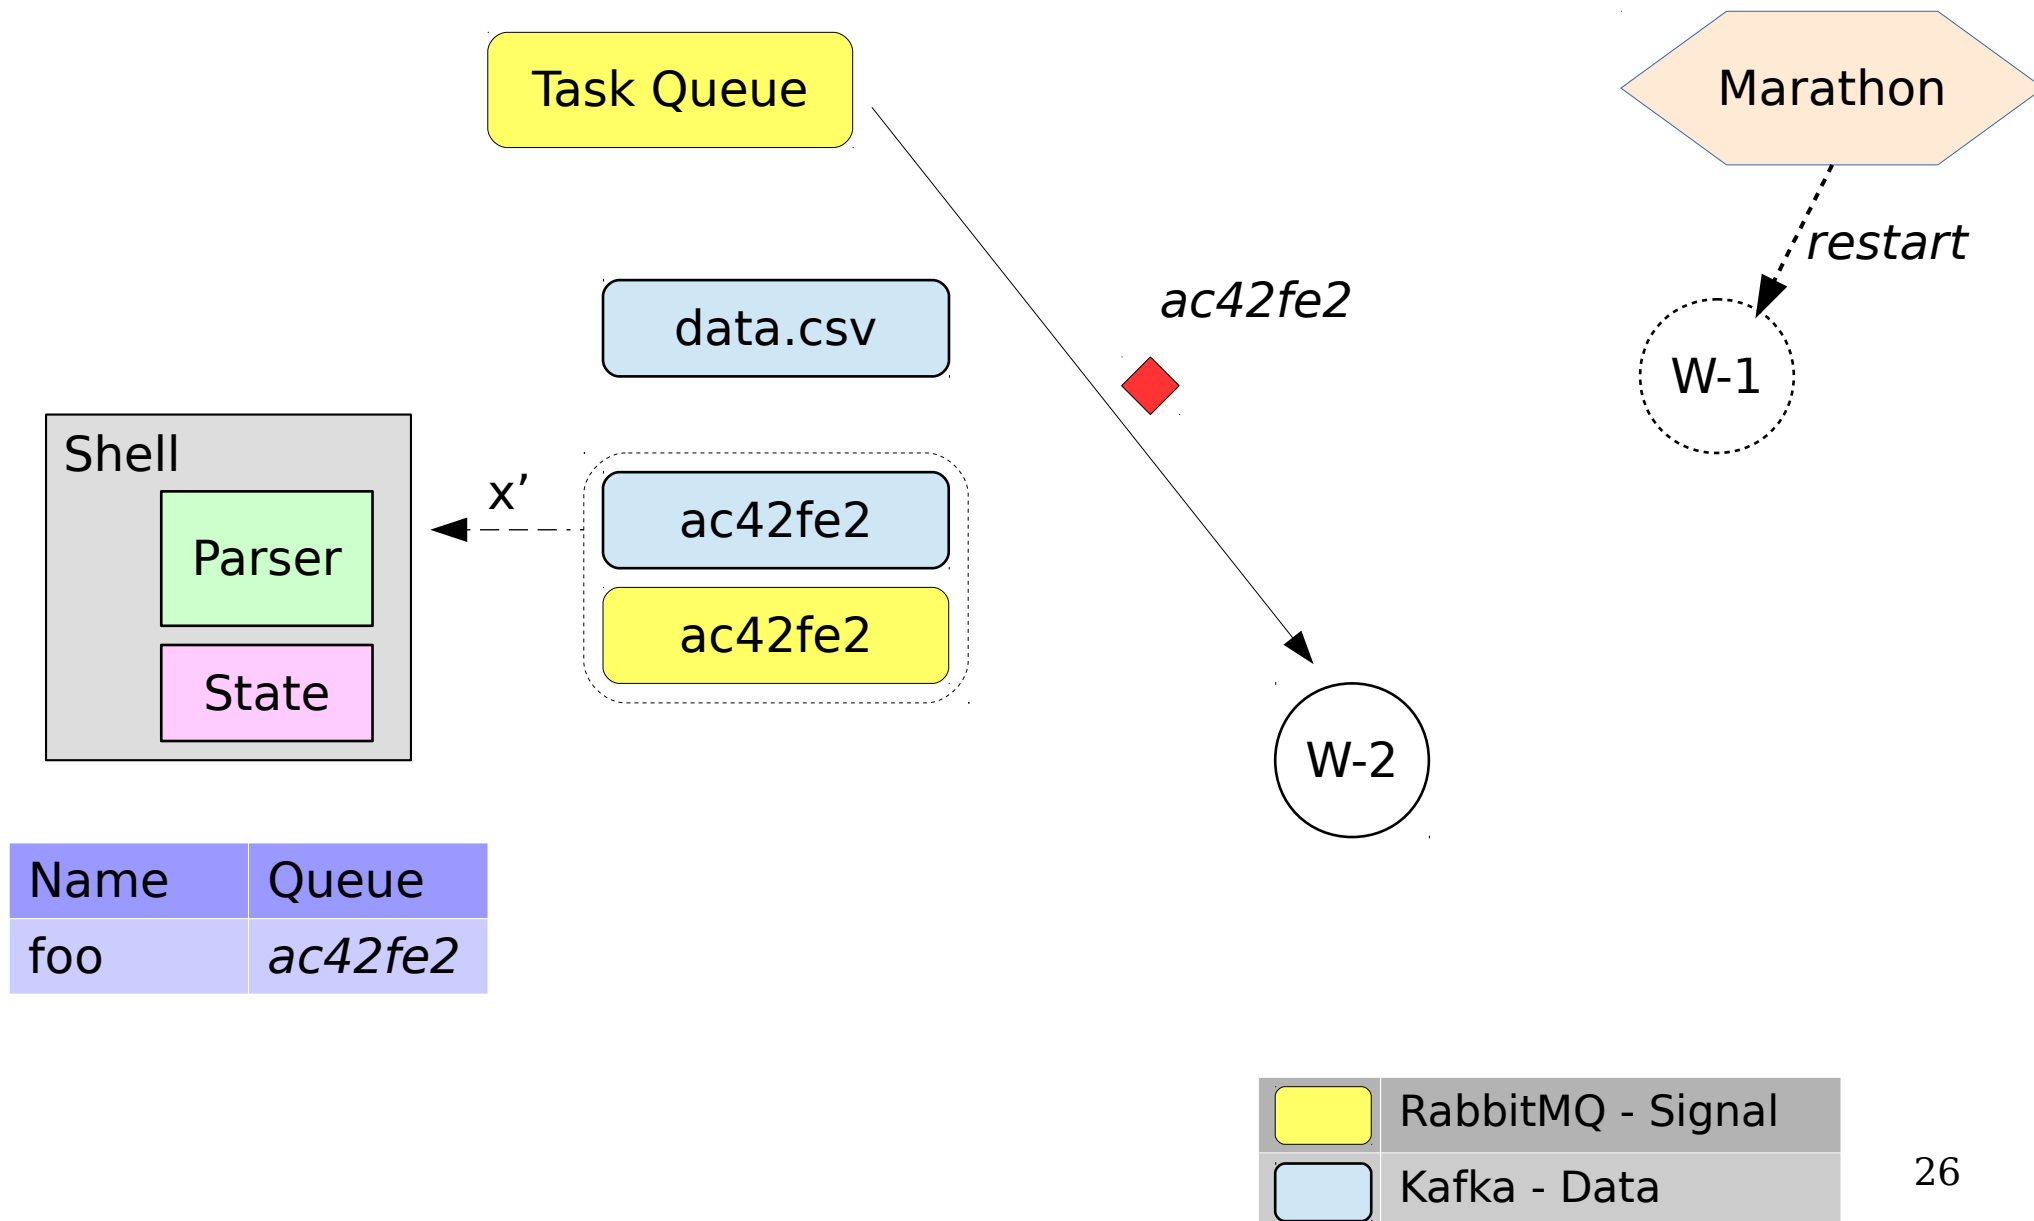
Marathon



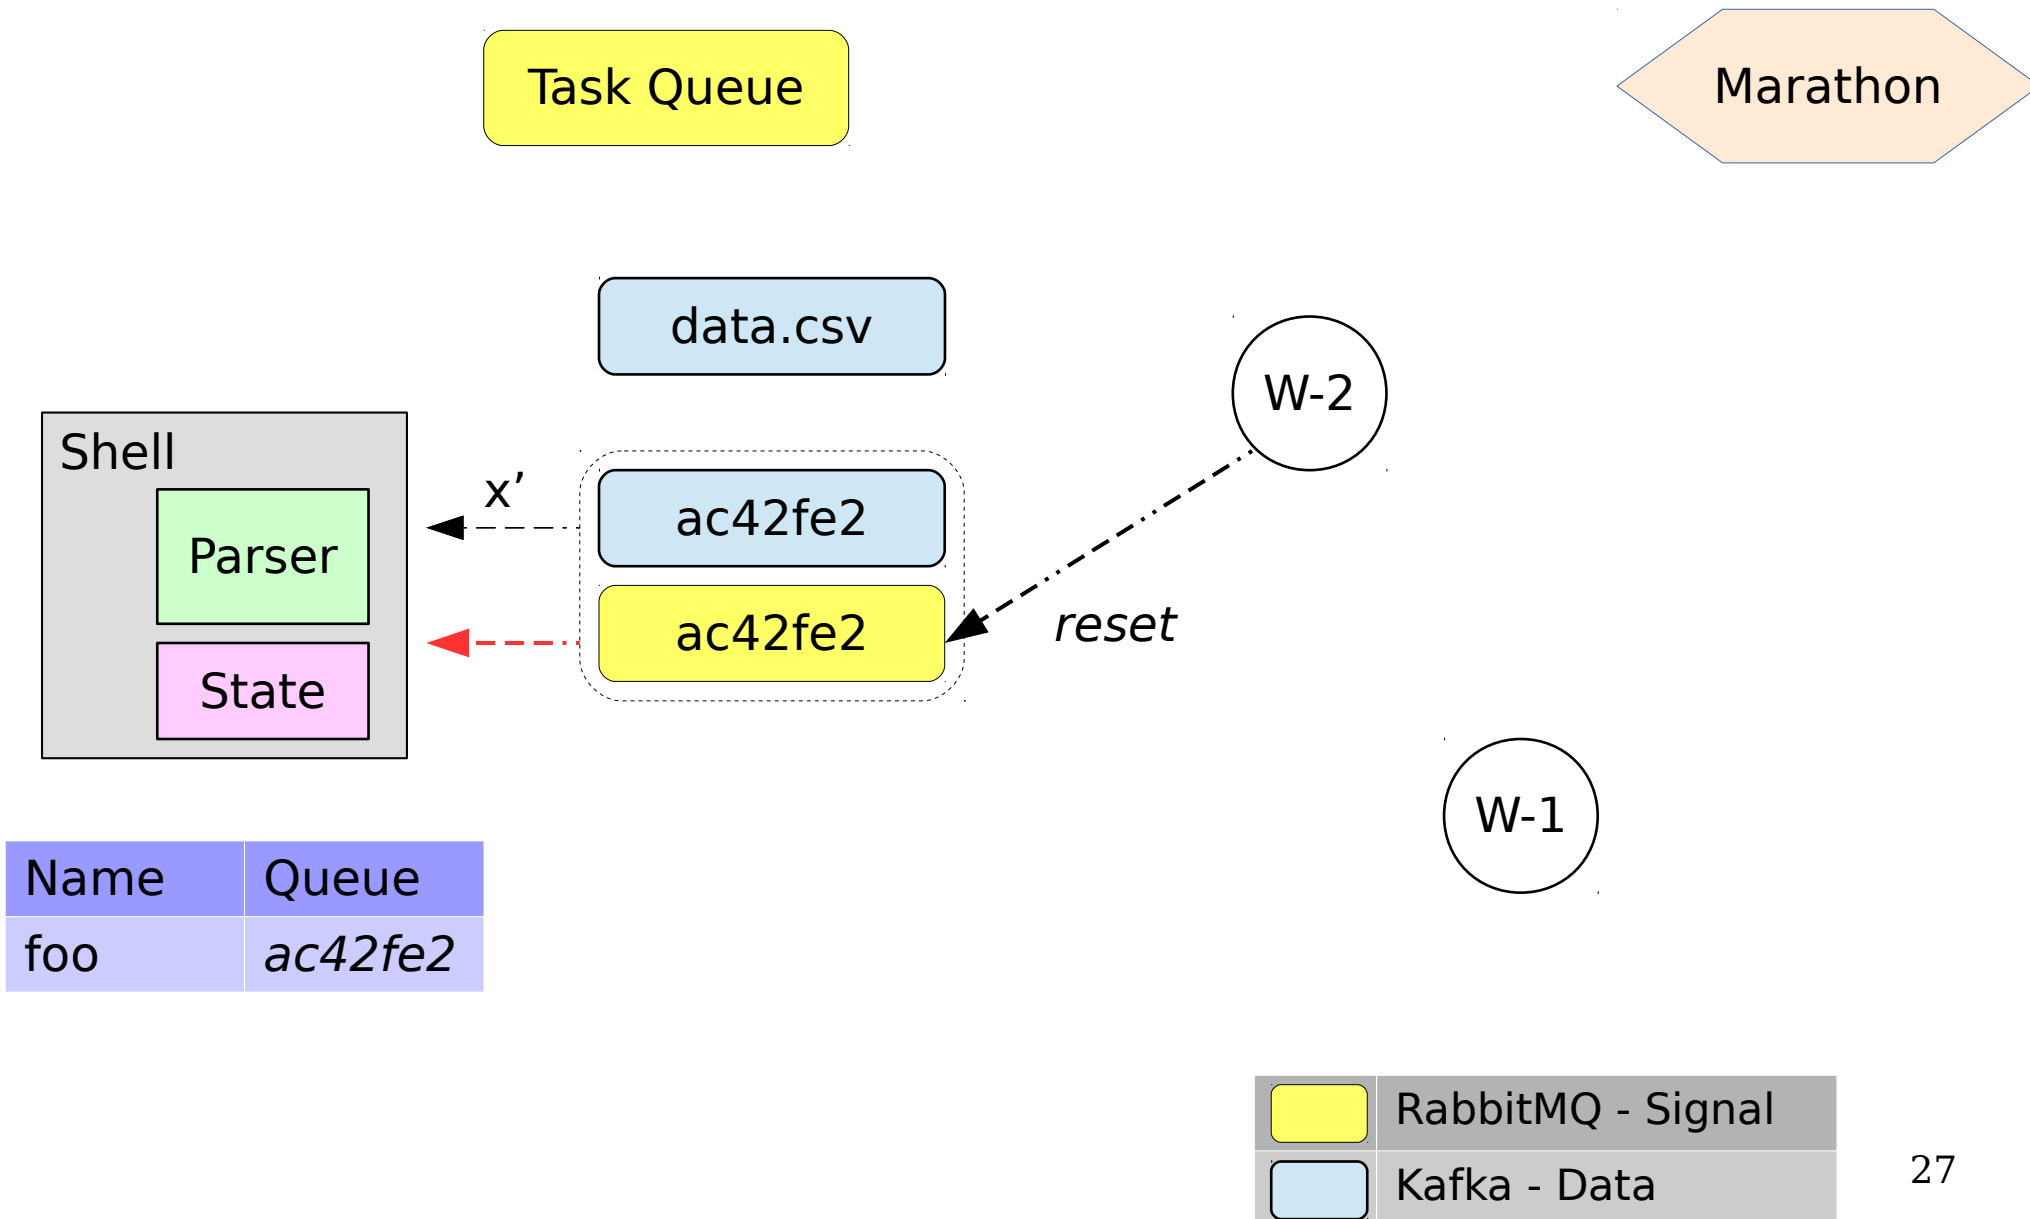
Name	Queue
foo	ac42fe2

	RabbitMQ - Signal
	Kafka - Data

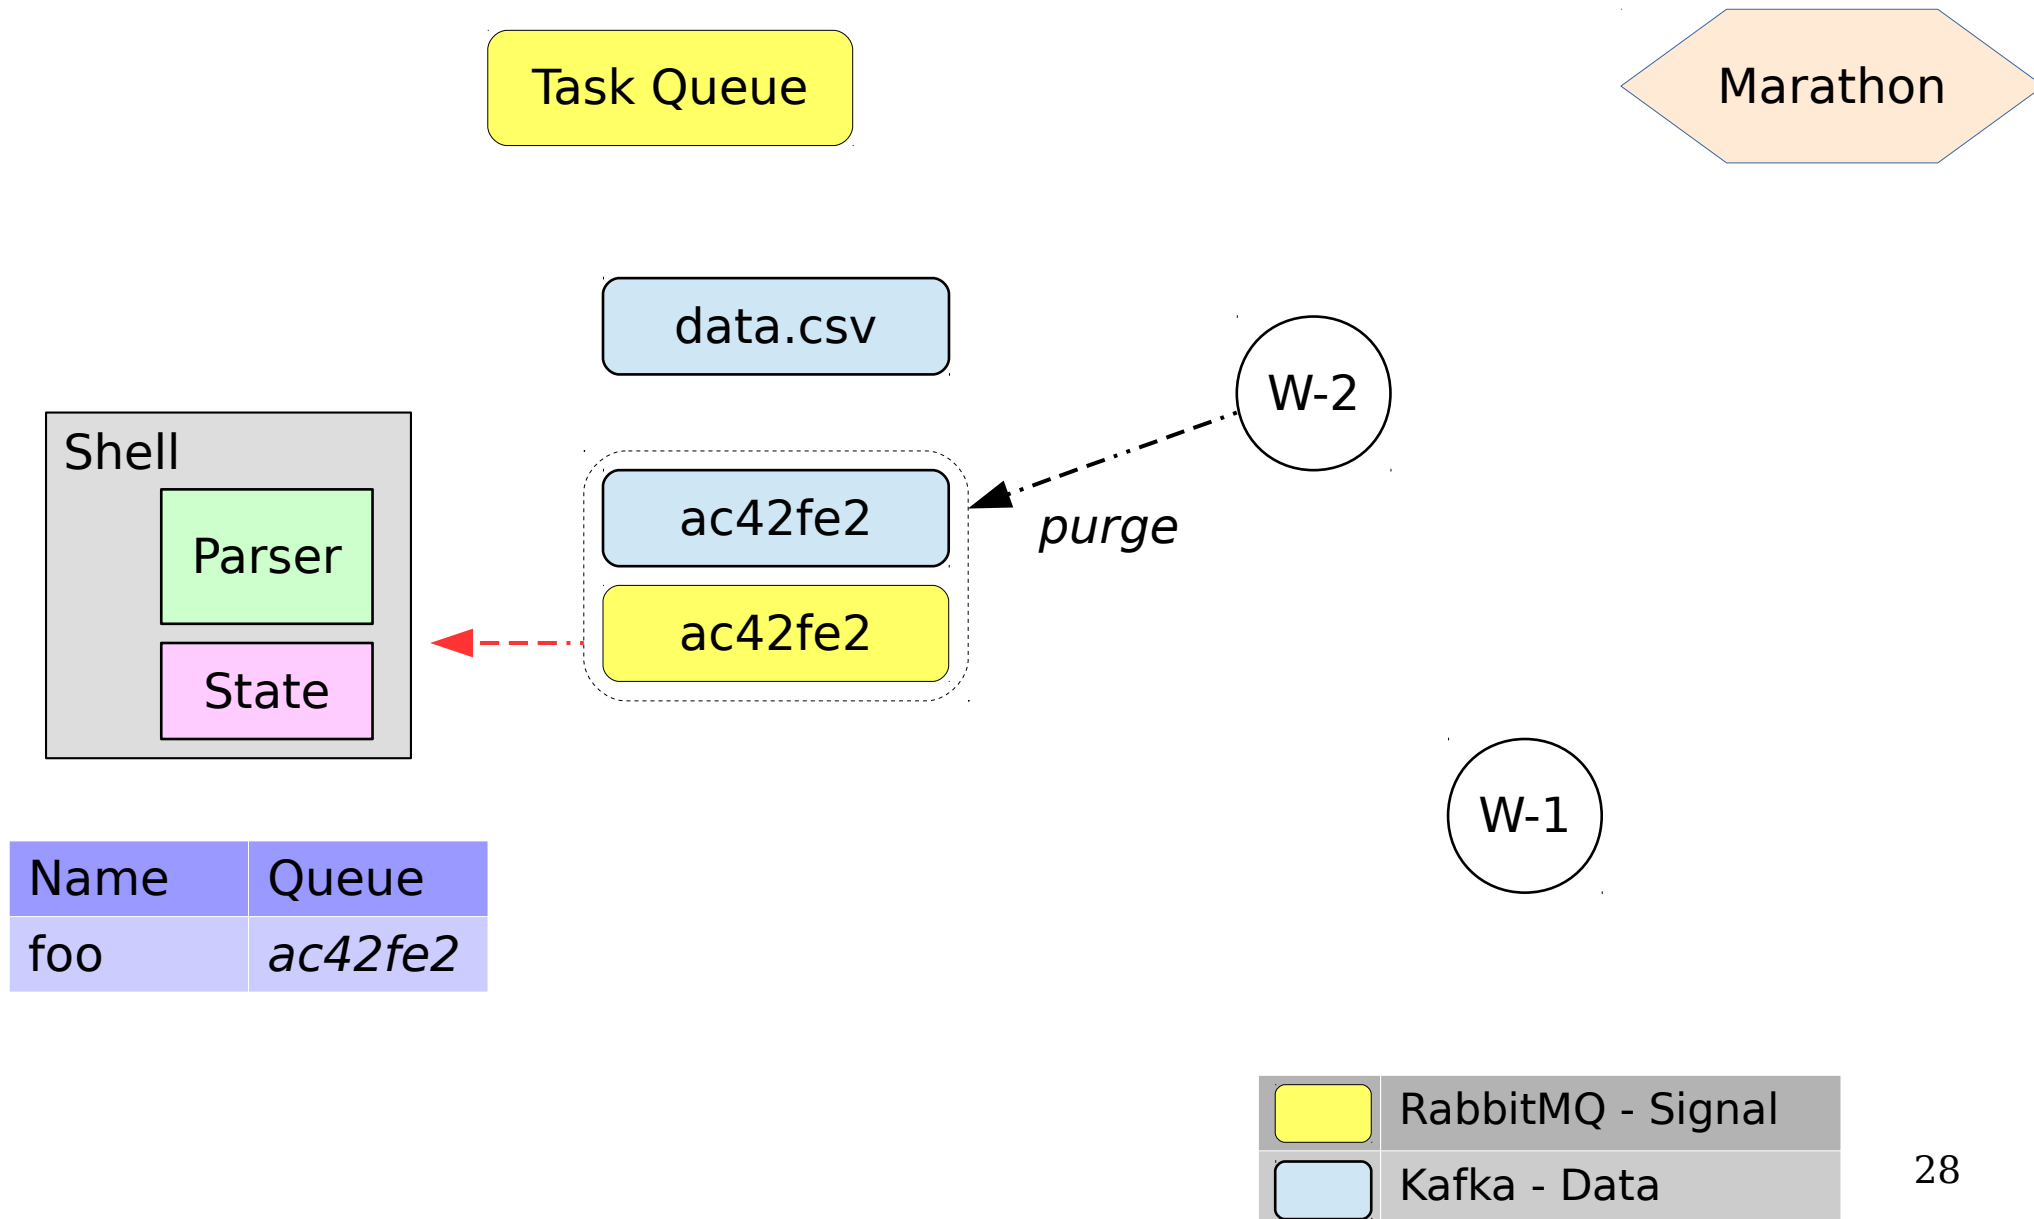
Drift Execution



Drift Execution



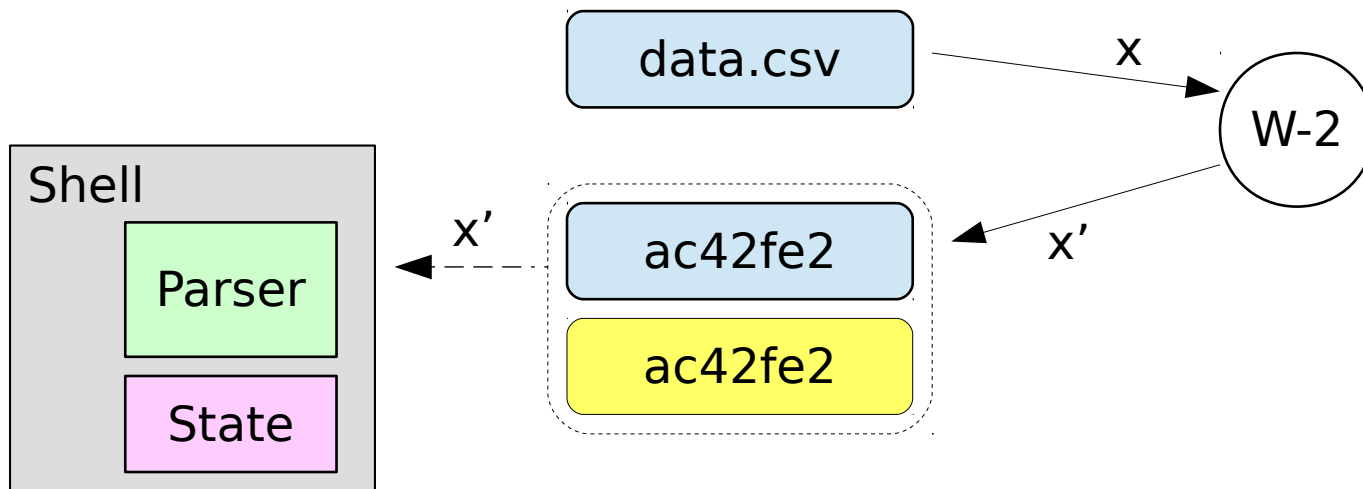
Drift Execution





Drift Execution

Task Queue

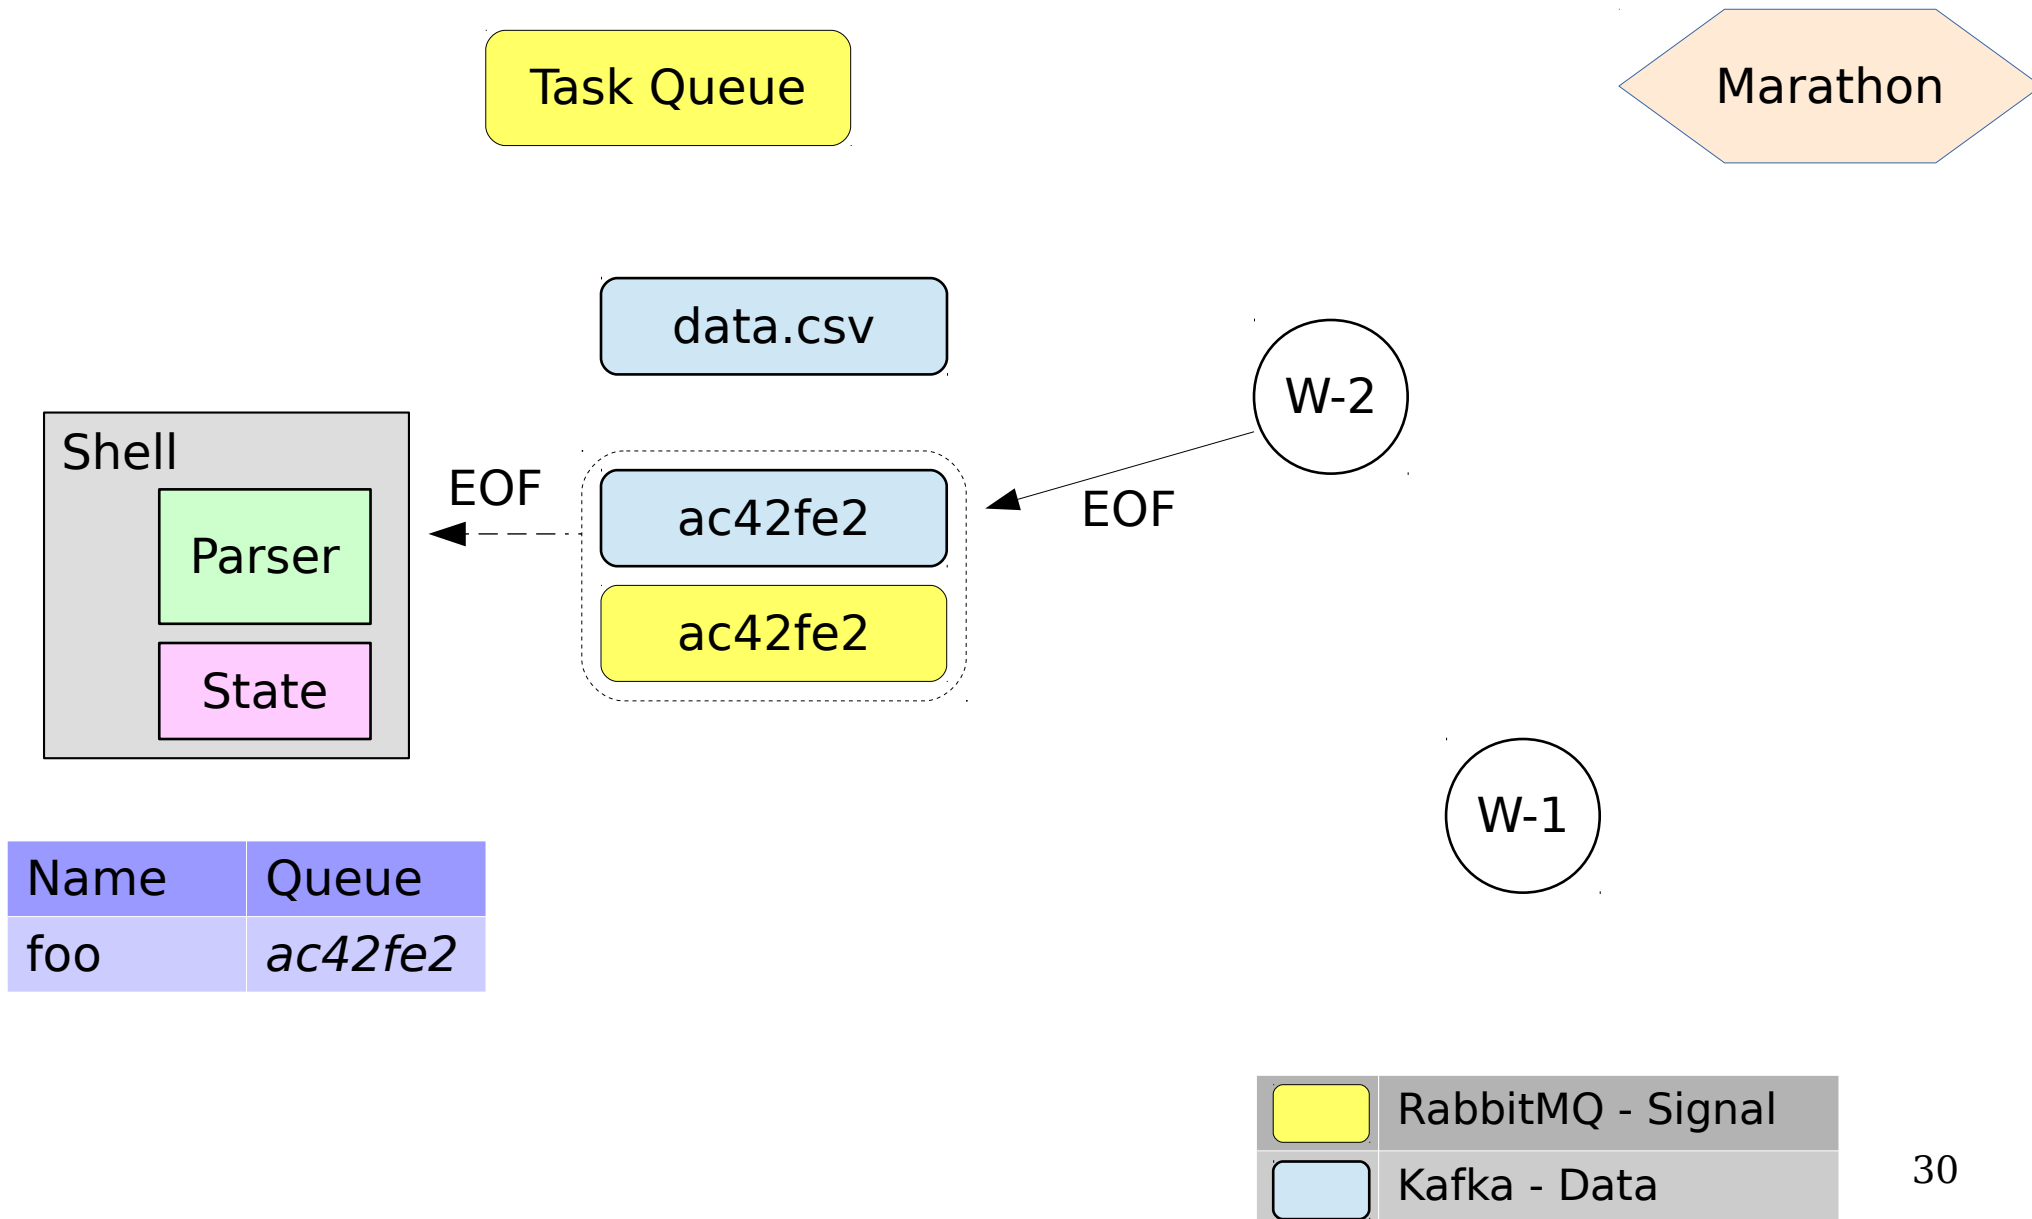
Marathon



Name	Queue
foo	ac42fe2

	RabbitMQ - Signal
	Kafka - Data

Drift Execution



Drift Execution

Task Queue

Marathon

data.csv

W-2

ac42fe2 ✓

ac42fe2



W-1

Shell

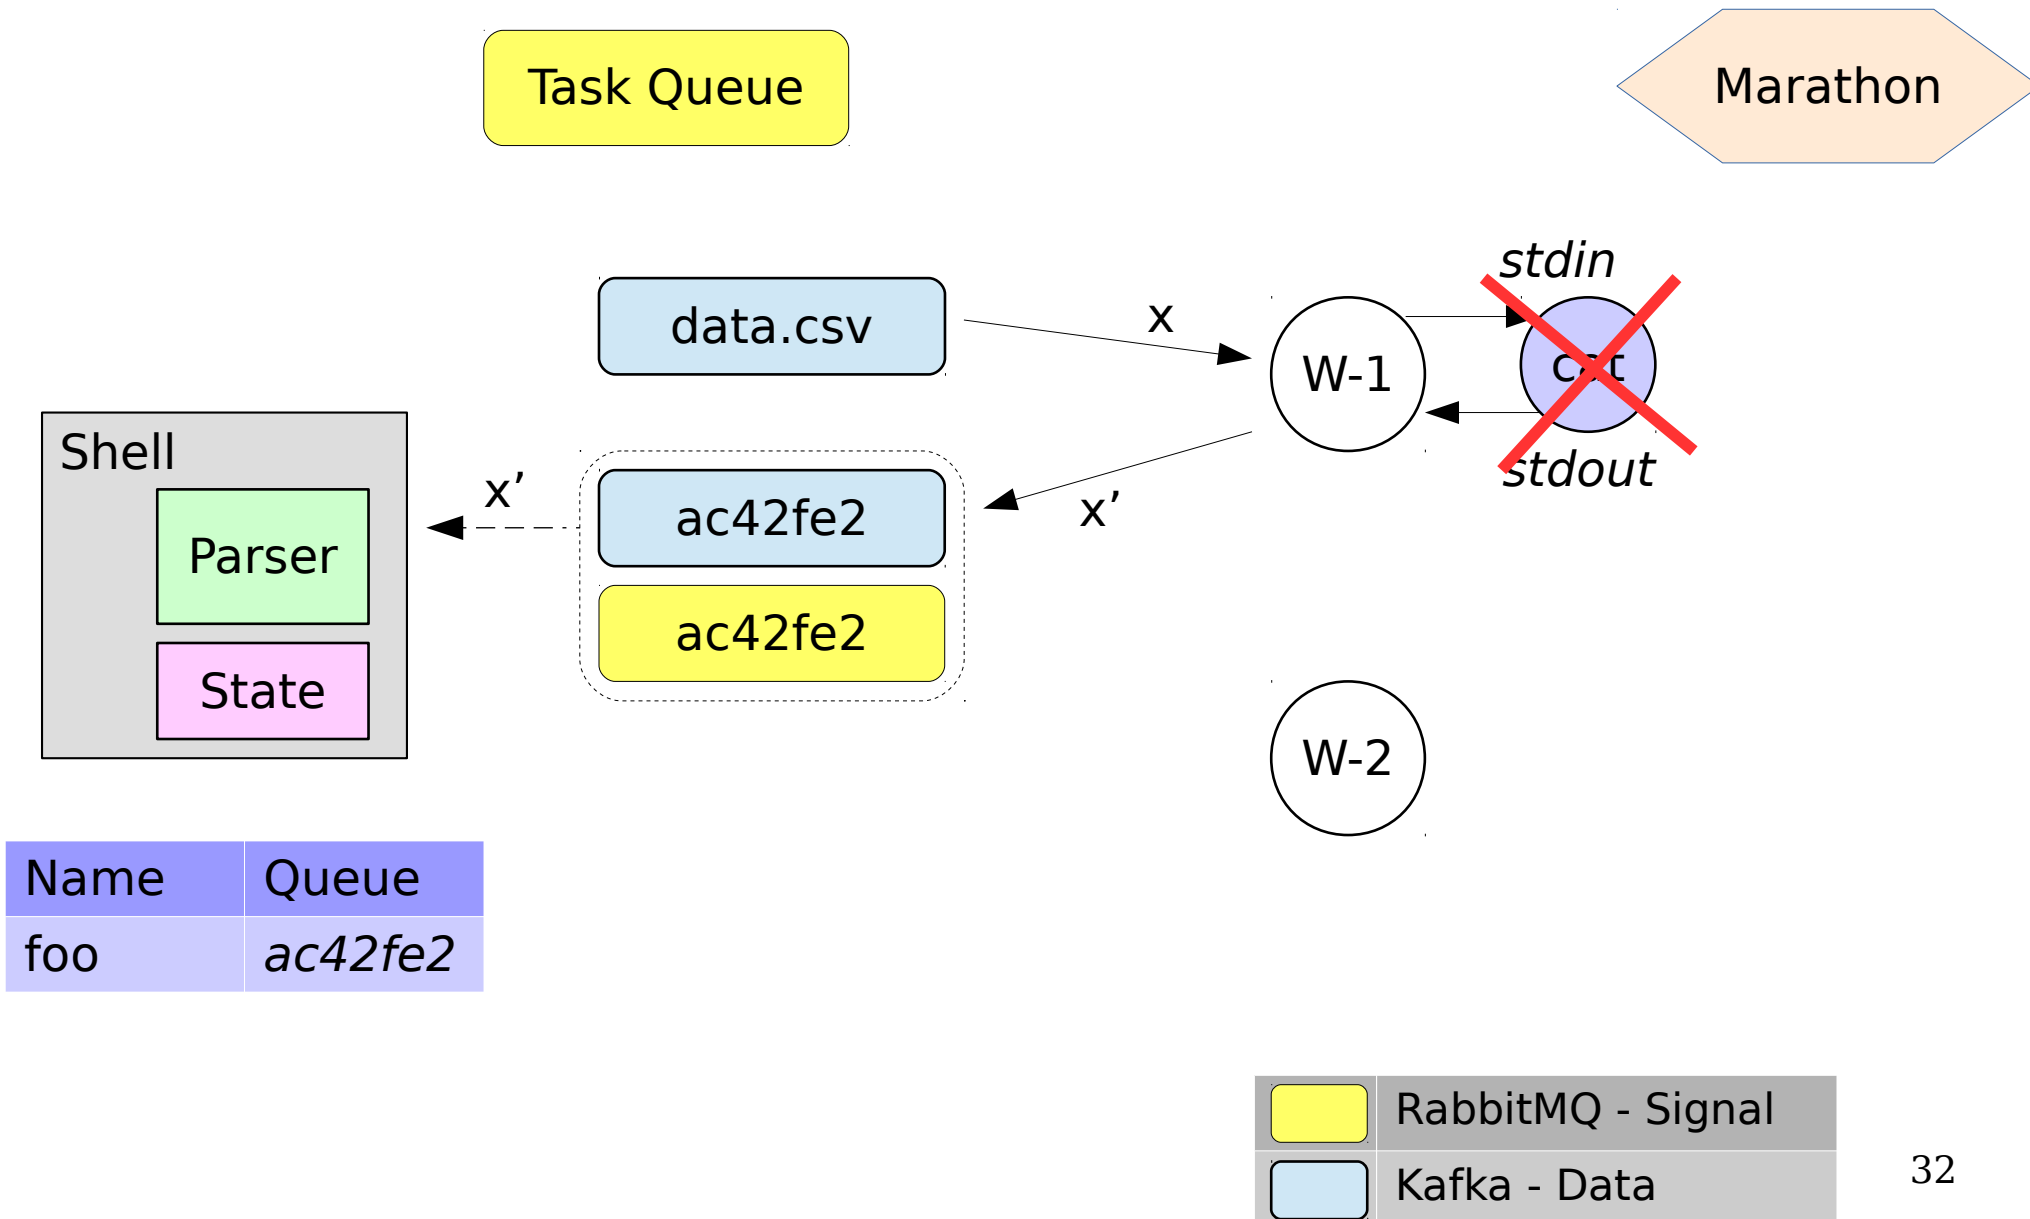
Parser

State

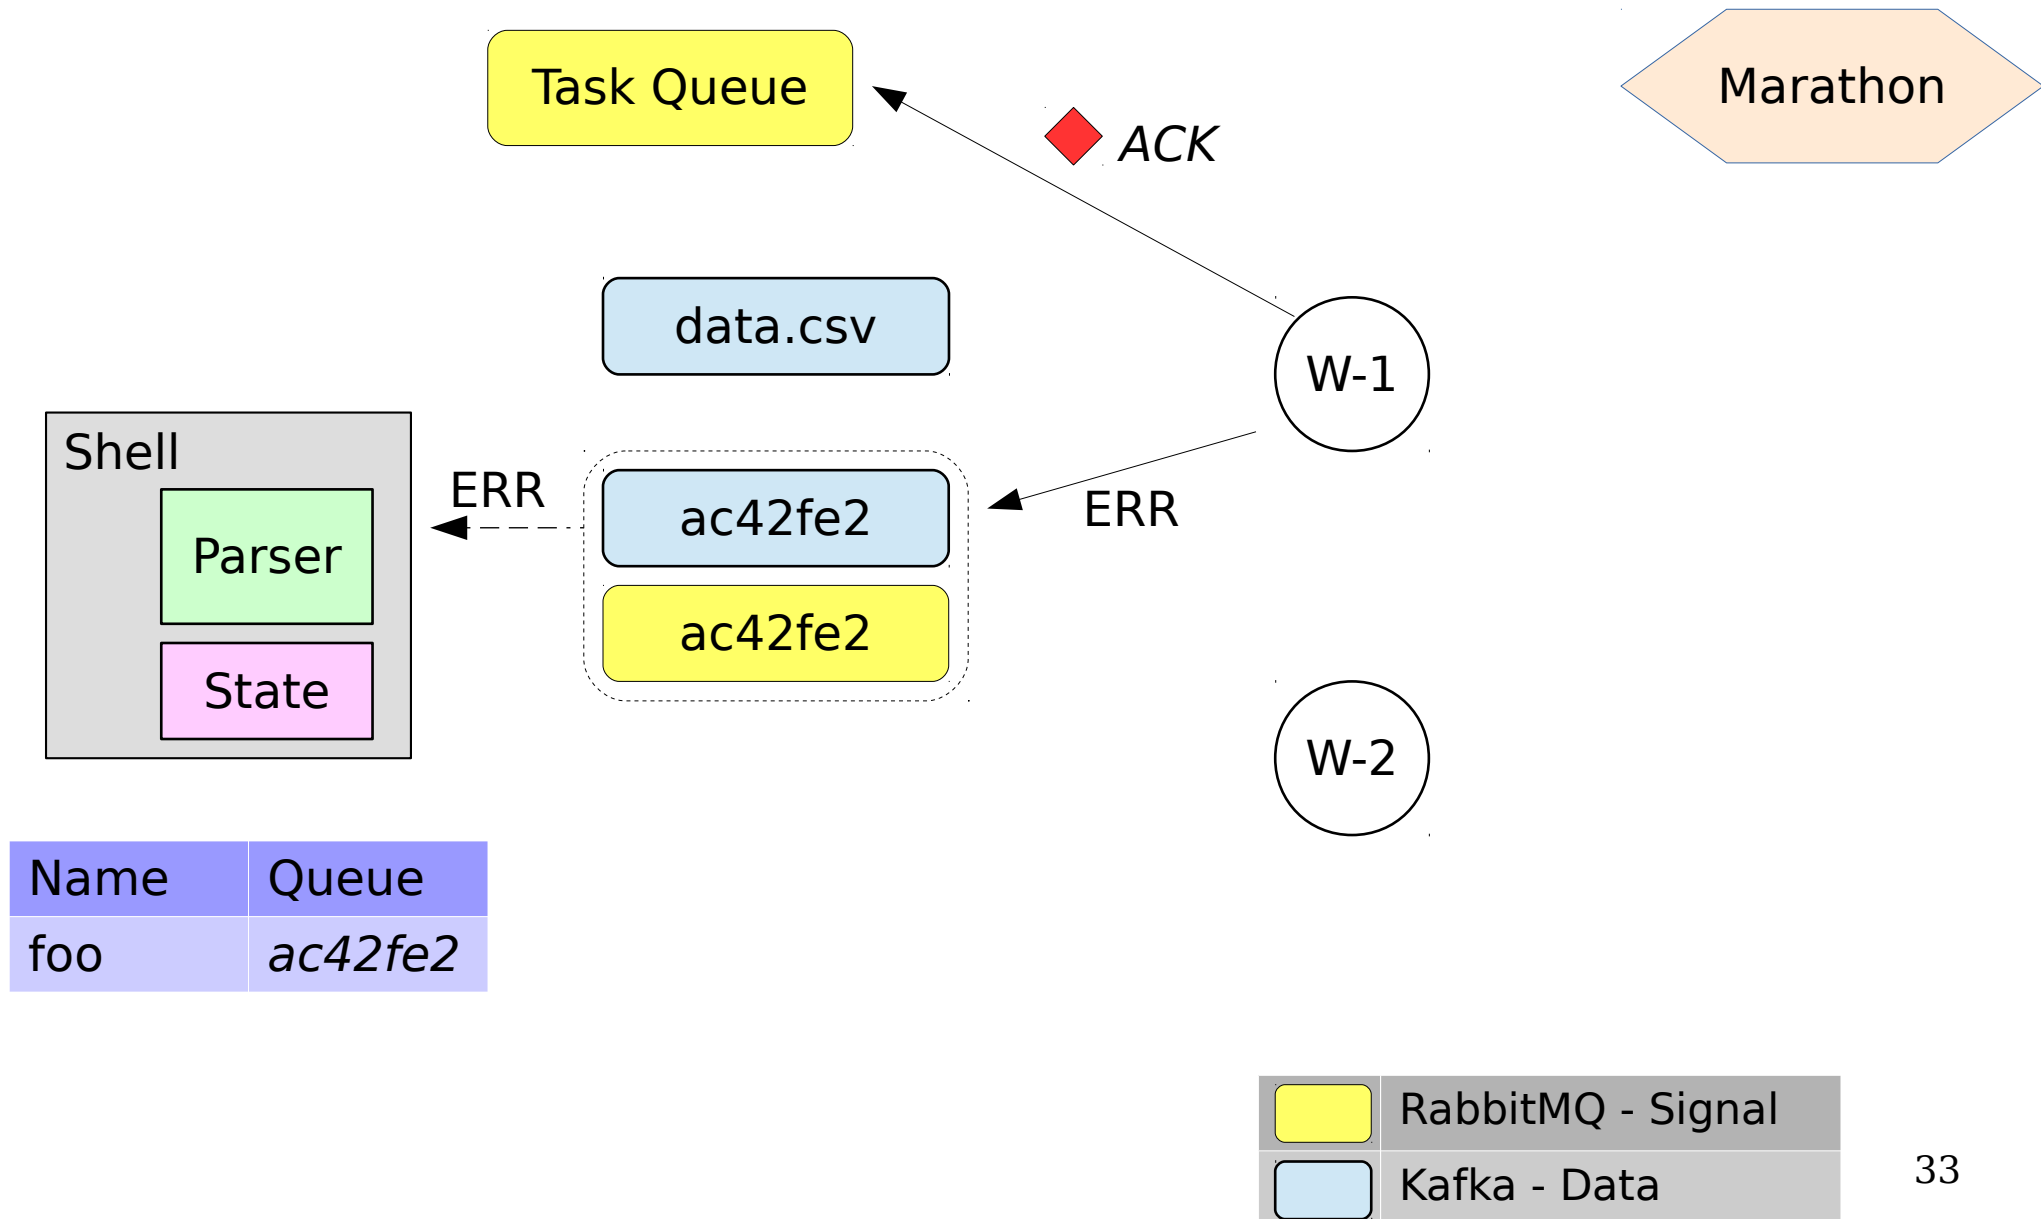
Name	Queue
foo	ac42fe2

	RabbitMQ - Signal
	Kafka - Data

Drift Execution



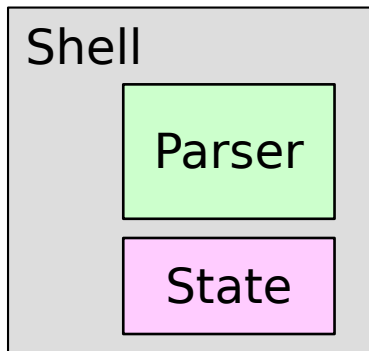
Drift Execution



Drift Execution

Task Queue

Marathon



data.csv

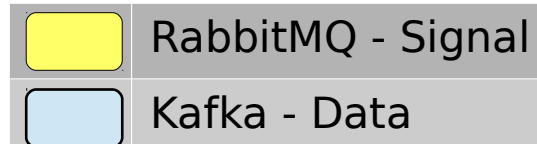
ac42fe2

ac42fe2

W-1

W-2

Name	Queue
foo	ac42fe2



Error Model

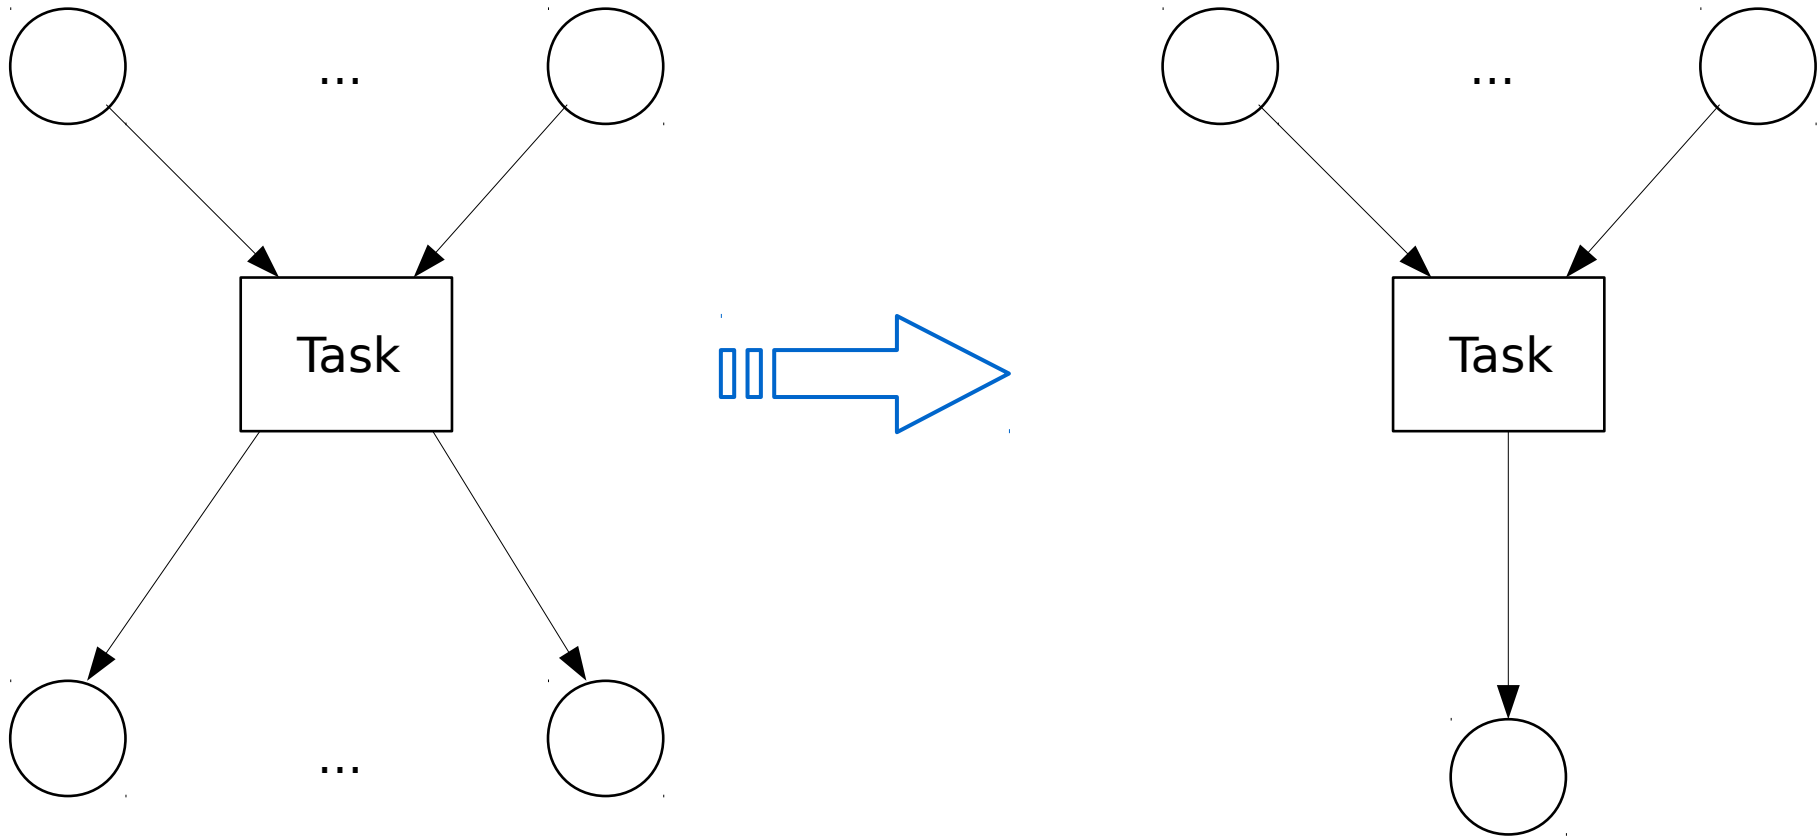
- Error Model: crash recovery
- System/Worker Failover → transparent
- Task Failure → permanent
 - because deterministic tasks
- queued tasks run until ERR token
- otherwise: task not accepted by interpreter

Invariants

System Invariants:

- Task scheduled exactly once or never
- Task executed exactly once or never
- Task always completed with EOF or ERR
- Task are deterministic
- possible to start unbounded tasks
 - language: unbounded
 - implementation: bounded

Invariants



Drift GUI

- Petri Net syntax natural fit for data + services
- different semantics:
 - no ‘occurrence rule’
 - no markings consumption
- BUT same properties like:
 - transition locality
 - async by default

Future Work

Major Problems:

- Semantics ...
- Types / Safety ...
- Platform ...

Minor Problems:

- Visualization
- Generalization / API